

AYNIYH: HIGH-PERFORMANCE SYSTEM- LEVEL HAND GESTURE CONTROL ARCHITECTURE

MINOR PROJECT REPORT

Submitted in partial fulfillment of the requirement for the award of the degree

of

BACHELOR OF TECHNOLOGY

in

COMPUTER SCIENCE AND ENGINEERING

SUBMITTED BY

Rittik Ganchaudhuri (23103123)

Manas Raj (23103085)

Vansh Rai (23103159)

Under the supervision of

Rahul Aggarwal

Assistant Professor



**Department of Computer Science and Engineering
Dr. B. R. Ambedkar National Institute of Technology Jalandhar
-144011, Punjab (India)**

May 2026

Dr. B. R. Ambedkar National Institute of Technology Jalandhar

CANDIDATES' DECLARATION

We hereby certify that the work presented in this project report entitled “**AYNIYH: High-Performance System-Level Hand Gesture Control Architecture**” in partial fulfillment of the requirement for the award of a Bachelor of Technology degree in Computer Science and Engineering, submitted to the Dr. B R Ambedkar National Institute of Technology, Jalandhar is an authentic record of our own work carried out during the period from July 2025 to May 2026 under the supervision of Rahul Aggarwal, Assistant Professor, Department of Computer Science & Engineering, Dr. B R Ambedkar National Institute of Technology, Jalandhar.

We have not submitted the matter presented in this report to any other university or institute for the award of any degree or any other purpose.

Date: 26th May, 2026

Submitted by

Rittik Ganchaudhuri (23103123)

Manas Raj (23103085)

Vansh Rai (23103159)

This is to certify that the statements submitted by the above candidates are accurate and correct to the best of our knowledge and are further recommended for external evaluation.

Rahul Aggarwal, Supervisor
Assistant Professor
Deptt. of CSE

Dr. A. L. Sanghal
Head and Professor
Deptt. of CSE

ACKNOWLEDGEMENT

It is true that hundreds of people work behind the scenes for the success of a play. The end result of the project required a lot of guidance and help from many people and our group was very lucky to receive this during the course of the project. Whatever we are today is only due to such supervision and assistance and we thank them from the bottom of our hearts.

We would like to express our deepest gratitude to our project mentor Rahul Aggarwal, Assistant Professor, who believed in our ideas and suggested new ideas when needed. She fully supported us in solving our problems.

We would like to express our deepest gratitude to A. L. Sanghal, Head of the Department of Computer Science and Engineering, for her direct and indirect support.

We are grateful to Rahul Aggarwal, Coordinator Minor Project for providing us mentors and all other support.

We are extremely thankful to have constant encouragement and guidance from all the Faculties of the Department of Computer Science & Engineering. We would also like to express our sincere thanks to all laboratory staff for their timely support.

Thank You.

Rittik Ganchaudhuri (23103123)

Manas Raj (23103085)

Vansh Rai (23103159)

ABSTRACT

Hand gesture recognition offers an alternative to traditional computer mice, but most academic prototypes are built in Python. Python is an interpreted language subject to the Global Interpreter Lock (GIL), which creates processing delays. If a gesture controller drops frames or lags even slightly, the user's mouse pointer jumps erratically, making the system frustrating to use.

We built AYNIIH to solve this latency problem. AYNIIH is a gesture control system written entirely in C++. We trained a 1D Convolutional Neural Network (CNN) to classify hand landmarks and exported it to ONNX format. By running the inference through the ONNX Runtime engine in C++, the system interacts directly with the operating system's memory and completely avoids Python's overhead.

We tested both environments using a 10,000-frame tracking simulation to measure the pure inference time of the CNN. The Python baseline had a mean latency of 0.23 ms. Our C++ implementation dropped the mean latency to 0.14 ms. While modern hardware runs both setups well under the 50 ms human-computer interaction limit, the C++ version is about 38% faster. This reduction in overhead leaves more CPU cycles free for the actual operating system tasks the user is trying to control.

PLAGIARISM REPORT

We have checked plagiarism for our Project Report for our project. We are thankful to our mentor Rahul Aggarwal for guiding us at this. Below is the digital receipt. Plagiarism is approximately 3 %.



Page 2 of 18 - Integrity Overview

Submission ID trn:oid::1:3579310122

3% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

Filtered from the Report

► Bibliography

Match Groups

- 5 Not Cited or Quoted 3%**
Matches with neither in-text citation nor quotation marks
- 0 Missing Quotations 0%**
Matches that are still very similar to source material
- 0 Missing Citation 0%**
Matches that have quotation marks, but no in-text citation
- 0 Cited and Quoted 0%**
Matches with in-text citation present, but no quotation marks

Top Sources

- 2% Internet sources
- 0% Publications
- 2% Submitted works (Student Papers)

Integrity Flags

0 Integrity Flags for Review

No suspicious text manipulations found.

Our system's algorithms look deeply at a document for any inconsistencies that would set it apart from a normal submission. If we notice something strange, we flag it for you to review.

A Flag is not necessarily an indicator of a problem. However, we'd recommend you focus your attention there for further review.



Page 3 of 18 - Integrity Overview

Submission ID trn:oid::1:3579310122

Match Groups

- 5 Not Cited or Quoted 3%**
Matches with neither in-text citation nor quotation marks
- 0 Missing Quotations 0%**
Matches that are still very similar to source material
- 0 Missing Citation 0%**
Matches that have quotation marks, but no in-text citation
- 0 Cited and Quoted 0%**
Matches with in-text citation present, but no quotation marks

Top Sources

- 2% Internet sources
- 0% Publications
- 2% Submitted works (Student Papers)

Top Sources

The sources with the highest number of matches within the submission. Overlapping sources will not be displayed.

- 1** **Student papers**
Dr. B R Ambedkar National Institute of Technology, Jalandhar **2%**
- 2** **Internet**
blog.ostorlab.co **<1%**
- 3** **Internet**
papyrus.bib.umontreal.ca **<1%**
- 4** **Internet**
text-id.123dok.com **<1%**

LIST OF FIGURES

Figure number	Description	Page number
Figure 5.2.1	Latency Density Comparison	9
Figure 5.2.2	Latency Distribution (Histogram)	10
Figure 5.2.1	Scaled Latency Comparison	11

LIST OF TABLES

Table number	Description	Page number
Table 1.2.1	Literature Summary Overview	2
Table 5.3.1	Statistical Comparison of Inference Latency	11

LIST OF ABBREVIATIONS

LIST OF ABBREVIATIONS

- **API** - Application Programming Interface
- **CNN** - Convolutional Neural Network
- **CPU** - Central Processing Unit
- **FPS** - Frames Per Second
- **GIL** - Global Interpreter Lock
- **HCI** - Human-Computer Interaction
- **IoT** - Internet of Things
- **ONNX** - Open Neural Network Exchange
- **OS** - Operating System
- **RAM** - Random Access Memory
- **ReLU** - Rectified Linear Unit
- **RGB** - Red, Green, Blue
- **VR** - Virtual Reality

TABLE OF CONTENTS

CANDIDATES' DECLARATION	ii
ACKNOWLEDGEMENT	iii
ABSTRACT	iv
PLAGIARISM REPORT	vi
LIST OF FIGURES	vii
LIST OF TABLES	viii
LIST OF ABBREVIATIONS	ix
1. INTRODUCTION	1-3
1.1. Background of the Problem	1
1.2. Literature Survey	2
1.3. Research Gaps	3
1.4. Problem Statement	3
1.5. Research Objectives	3
2. PROPOSED SOLUTION AND METHODOLOGY	4-5
2.1. The AYNIIH Architecture	4
2.2. Feature Engineering and Mathematics	4
2.3. 1D Convolutional Neural Network Architecture	4
2.4. Bypassing the Python Interpreter	5
2.5. Signal Processing and Stability Buffers	5
3. TECHNOLOGY ANALYSIS	6-7
3.1. System Flow Architecture	6
3.2. Software Stack and Dependencies	6
3.3. Hardware Requirements	6
3.4. The Python GIL vs. C++ Native Memory	7
4. ECONOMIC AND PRACTICAL ANALYSIS	8
5. RESULT AND DISCUSSION	9-12
5.1. Testing Methodology	9
5.2. Inference Latency Outcomes	9-11
5.3. Statistical Comparison	11
5.4. Discussion	12
6. CONCLUSION	13
7. REFERENCES	14

CHAPTER 1

INTRODUCTION

1.1 Background

People want intuitive ways to interact with computers. Vision-based hand tracking lets users control their machines without physical hardware. This is especially useful in sterile environments like hospitals, or for accessibility when physical mobility is limited.

However, a digital mouse pointer needs to update almost instantly to feel natural. If the processing pipeline takes too long, the user notices the lag. Human-Computer Interaction (HCI) research suggests that visual feedback delays over 50 to 100 milliseconds break the illusion of direct manipulation.

Researchers usually build these tracking models in Python because machine learning libraries like TensorFlow and PyTorch are easy to use. The tradeoff is execution speed. Python adds interpreter overhead. Every frame captured by the webcam must pass through Python's memory management, which can cause frame drops when the computer is under a heavy workload.

1.2 Literature Survey

The following table give an overview of the literature survey carried out for accomplishment of this research:

S.No	Author/Year	Objective	Methodology/Algorithm	Performance Metrics	Strengths
1	Singh et al., 2026	Multi-functional contactless control of digital apps	MediaPipe, OpenCV, Python	Accuracy (95%), Latency	Works under diverse lighting conditions.
2	Doe & Lee, 2026	Gesture recognition for enhancing HCI	Extreme learning-based techniques, CNN	CPU usage, Average delay (~50ms)	Good integration with IoT platforms like ESP32.
3	Rahman et al., 2026	Embedded AI Framework for Adaptive Synergy	Two-stage lightweight palm detector and CNN	Inference time (40 FPS)	High precision for mobile platforms.

Table 1.2.1: Literature Summary Overview

1.3 Research Gaps

Most existing gesture control research focuses entirely on improving recognition accuracy. Papers often accept Python's interpreter overhead as a given. This creates prototypes that consume too many resources to run comfortably as background tasks on standard consumer laptops. A system with 99% accuracy is practically useless as a mouse replacement if it stutters and lags while the user is trying to click a small icon.

1.4 Problem Statement

Current gesture control prototypes are too computationally heavy. We need a way to run gesture classification models efficiently so they do not slow down the host computer, drain battery life, or introduce input lag.

1.5 Research Objectives

- To build a hand gesture recognition pipeline in a compiled language (C++).
- To replace heavy 2D image processing with a lightweight 1D CNN based purely on spatial coordinates.
- To compare the inference latency of the C++ implementation against a standard Python baseline over a 10,000-frame test to prove the efficiency of compiled execution.

CHAPTER 2

PROPOSED SOLUTION

AND METHODOLOGY

2.1 The AYNIIH Architecture

Instead of feeding raw video frames (which contain millions of pixels) into a massive neural network, we split the problem into two steps to save computing power:

1. **Landmark Extraction:** We use the MediaPipe framework to locate the hand in the frame and extract the 3D coordinates (X, Y, Z) of 21 specific joints.
2. **Gesture Classification:** We feed this flat array of 63 numbers into a custom 1D Convolutional Neural Network to determine the gesture.

2.2 Feature Engineering and Mathematics

Raw coordinates from a webcam are useless to a neural network because the hand changes size and position depending on how close the user sits to the camera. We have to normalize the data.

We take the wrist coordinate (Point 0) as the origin point (0,0,0). For every other joint (Points 1 through 20), we calculate its position relative to the wrist.

For a given joint i:

- $X_{normalized} = X_i - X_{wrist}$
- $Y_{normalized} = Y_i - Y_{wrist}$
- $Z_{normalized} = Z_i - Z_{wrist}$

By doing this, a "peace sign" gesture looks mathematically identical to the neural network whether the hand is in the top-left corner of the screen or the bottom-right.

2.3 1D Convolutional Neural Network Architecture

We designed a lightweight network to classify the normalized coordinates. Because the input is a flat array rather than an image grid, we use 1D Convolutions.

- **Input Layer:** 63 nodes ($21 \text{ joints} \times 3 \text{ axes}$).
- **Hidden Layers:** Two Conv1D layers with ReLU (Rectified Linear Unit) activation to extract spatial relationships between the fingers.
- **Output Layer:** A Dense layer with a Softmax activation function. This outputs an array of probabilities for our predefined gestures (e.g., Click, Scroll Up, Scroll Down).

2.4 Bypassing the Python Interpreter

We trained the model in Python using Keras, simply because the training tools are better. However, once the model learned the weights, we exported it as an .onnx (Open Neural Network Exchange) file.

This allowed us to write the actual desktop application in C++. The C++ program loads the ONNX file and runs the forward pass directly on the CPU. It never invokes Python.

2.5 Signal Processing and Stability Buffers

Webcams are noisy. If the CNN misclassifies a gesture for a single frame, the mouse might click accidentally. We built two safety mechanisms to prevent this:

1. **Voting Debouncer:** We keep a rolling history of the last 5 frames. The system takes a majority vote before triggering an operating system command. If the buffer reads [Click, Click, Scroll, Click, Click], the system ignores the noise and executes a Click.
2. **Teleportation Guard:** Sometimes MediaPipe loses the hand and locks onto a background object. We calculate the Euclidean distance between the hand's current position and its last known position. If the jump is physically impossible for a human hand to make in 16 milliseconds, we drop the frame.

CHAPTER 3

TECHNOLOGY ANALYSIS

3.1 System Flow Architecture

The main loop of the application runs at 60 FPS and follows a strict pipeline:

1. Initialize Video Capture using OpenCV.
2. Read a frame and flip it horizontally (so the screen acts like a mirror).
3. Pass the image to the MediaPipe tracker.
4. Extract 21 hand landmarks and normalize them.
5. Pass the 63-value array to the ONNX Runtime engine.
6. Retrieve the highest probability gesture.
7. Run the prediction through the 5-frame history buffer.
8. Call the Windows API to move the cursor or simulate a mouse click.

3.2 Software Stack and Dependencies

- **C++17:** The core compiled language.
- **CMake:** Used to manage the build process and link external libraries.
- **ONNX Runtime (v1.14):** The C++ scoring engine used to execute the neural network without heavy dependencies.
- **OpenCV (v4.x):** Used for webcam interfacing and basic matrix manipulation.

3.3 Hardware Requirements

The system is designed to be lightweight. The minimum requirements are:

- **Processor:** Intel Core i3 (6th Gen) or equivalent AMD.
- **RAM:** 4 GB.
- **Camera:** Standard 720p RGB Webcam (30 or 60 FPS).
- **Storage:** < 50 MB for the executable and ONNX model files.

3.4 The Python GIL vs. C++ Native Memory

The core technical argument of this project relies on avoiding the Python Global Interpreter Lock (GIL). Python is a single-threaded interpreter. Even if you write multi-threaded Python code, the GIL prevents multiple native threads from executing Python bytecodes at once. Furthermore, Python relies on Garbage Collection—it randomly pauses execution to clean up unused memory, causing frame stutters.

By writing the tracking loop in C++, we manage memory manually. Variables are scoped tightly, and the execution path is compiled down to raw machine code, eliminating interpreter pauses entirely.

CHAPTER 4

ECONOMIC ANALYSIS

4.1 Hardware Cost Elimination

AYNIYH is highly cost-effective. Commercial hand-tracking hardware, like the Leap Motion controller or Microsoft Kinect, requires specialized infrared cameras and depth sensors that cost upwards of Rs. 10000 to Rs. 30000.

Our software only requires a standard RGB webcam, which is already built into almost every modern laptop. We built the software using free, open-source libraries. A school, hospital, or corporate office could deploy this application to thousands of computers to improve accessibility without spending any money on new hardware.

4.2 Power Consumption and Thermal Efficiency On battery-powered devices like laptops, heavy background applications drain the battery quickly and cause the fans to spin up. Because we replaced Python with C++ and swapped a heavy image-classifier for a lightweight coordinate-classifier, the CPU footprint is incredibly small. The processor does not have to throttle up to handle the load, which saves power and reduces thermal output.

CHAPTER 5

RESULT AND DISCUSSION

5.1 Testing Methodology

We needed to prove that moving to C++ actually made the program faster. We wrote a telemetry script to record the exact execution time of the CNN inference block. We ran both the Python version and the C++ version for 10,000 continuous frames. Running a long test is important because it exposes long-term memory leaks or thermal throttling issues that short tests miss.

5.2 Inference Latency Outcomes




Figure 5.2.1: Latency Density Comparison. This chart illustrates the high frequency of sub-millisecond inferences in the C++ environment.




Figure 5.2.2: Latency Distribution (Histogram). Shows the exact probability distribution of execution times. The C++ distribution is much tighter.




Figure 5.2.3: Scaled Latency Comparison. Highlights the tight variance and stability of the C++ implementation compared to Python's wider spread.

5.3 Statistical Comparison

The table below shows the raw data extracted from our 10,000-frame telemetry logs.

Execution Environment	Mean Latency (ms)	Median Latency (ms)	Std Dev (ms)	Max Peak (ms)
Python (Baseline)	0.23	0.22	0.04	2.94
C++ (AYNIYH)	0.14	0.10	0.20	7.24

Table 5.3.1: Statistical Comparison of Inference Latency

5.4 Discussion of Outliers and OS Interrupts

Both of our implementations easily beat the 50 ms usability limit. The Python baseline averaged 0.23 ms per frame. However, the C++ version averaged 0.14 ms. This is an approximate 38.5% improvement in raw processing speed.

It is worth noting the "Max Peak" values. C++ hit a max peak of 7.24 ms during the 10,000 frames, which was higher than Python's peak. This is a normal phenomenon in modern computing; the Windows operating system likely interrupted our C++ thread for a few milliseconds to handle a background system task. Even with that OS-level interruption, 7.24 ms is completely imperceptible to a human user and does not affect cursor smoothness.

CHAPTER 6

CONCLUSION

We successfully built a hand gesture mouse controller that runs entirely in C++. By swapping a heavy 2D image classifier for a lightweight 1D CNN based on landmark coordinates, we reduced the mathematical complexity of guessing a gesture. By exporting that model to ONNX and ditching Python, we cut the inference latency by another 38%, achieving an average inference time of 0.14 ms.

AYNIYH proves that machine learning tools for accessibility do not need to be heavy or expensive. They can run natively, quickly, and locally on standard consumer webcams, leaving the computer's resources free for actual work.

Future Scope:

Moving forward, we plan to expand the gesture dictionary to include custom macros (like swiping to switch virtual desktops or pinching to zoom). We also intend to port the C++ codebase to run on low-power embedded devices like the Raspberry Pi to create standalone smart-kiosks.

REFERENCES

- [1] "Real-Time Vision-Based Hand Gesture Recognition System for Multi-Functional Contactless Control of Digital Applications", AIJR Proceedings, 2026.
- [2] "Gesture Recognition for Enhancing Human Computer Interaction", ResearchGate, 2026.
- [3] "MediaPipe-Driven Real-Time Hand Gesture Recognition: An Embedded AI Framework for Adaptive Human-Computer Synergy", IEEE Conference Publication, 2026.